

Προγραμματισμός σε C++ & Python & Εφαρμογές στη Ναυπηγική & Ναυτική Μηχανολογία

ΣΝΜΜ 2019

Μάθημα 6B: Εφαρμογή: Neural Networks. Εφαρμογή:
Hardware. Πλατφόρμες. Πρωτόκολλα. Βασικό I/O

Γεώργιος Παπαλάμπρου
Επίκουρος Καθηγητής ΕΜΠ
george.papalambrou@lme.ntua.gr

Εργαστήριο Ναυτικής Μηχανολογίας (Κτίριο Α)
Σχολή Ναυπηγών Μηχανολόγων Μηχανικών
Εθνικό Μετσοβιό Πολυτεχνείο

April 11, 2019

Περιεχόμενα

- 1 Εφαρμογή: Neural Networks. Machine Learning
- 2 Εφαρμογή: Hardware. Πλατφόρμες. Πρωτόκολλα. Βασικό I/O

Περιεχόμενο Μαθήματος

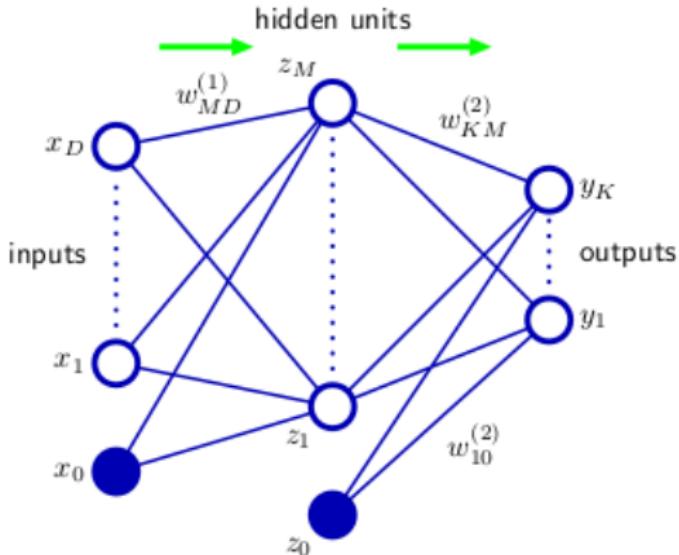
- Εβδομάδα 1. A. Εισαγωγή. Η γλώσσα. Το περιβάλλον Linux. Command line. Python interpreter. Ιστοσελίδα μαθήματος. Βιβλιογραφία. Editors: Sublime, Spyder. B. Εισαγωγή στην γλώσσα Python. Hello World.
- Εβδομάδα 2. A. Data types. Loops. Control. B. Παραδείγματα
- Εβδομάδα 3. Functions. Modules
- Εβδομάδα 4. OOP. Classes
- Εβδομάδα 5. A. Παραδείγματα: Μέτρηση και επεξεργασία δεδομένων. B. Errors-Exceptions.
- Εβδομάδα 6.
 - A. Βιβλιοθήκες Numpy, SciPy. Γραφικά
 - B. Εφαρμογή: Neural Networks. Machine Learning. Εφαρμογή: Hardware. Πλατφόρμες. Πρωτόκολλα. Βασικό I/O

Βασικές Πηγές ΝΔ

- Neural Networks: A Comprehensive Foundation, Simon Haykin, Macmillan, 1994. [Ch. 1,2,4]
- Pattern Recognition and Machine Learning, Christopher Bishop, Springer, 2006. [Ch. 5]
[Διαθέσιμο pdf στο: <https://www.microsoft.com/en-us/research/publication/pattern-recognition-machine-learning/>]
- Machine Learning, Tom Mitchell, McGraw Hill, 1997. [Ch. 4]
- Course MIT 6.S191: Introduction to Deep Learning, www:introtodeeplearning.com, [Slides Lesson 1]

To νευρωνικό δίκτυο

Figure 5.1 Network diagram for the two-layer neural network corresponding to (5.7). The input, hidden, and output variables are represented by nodes, and the weight parameters are represented by links between the nodes, in which the bias parameters are denoted by links coming from additional input and hidden variables x_0 and z_0 . Arrows denote the direction of information flow through the network during forward propagation.



[Pattern Recognition and Machine Learning, Bishop]

To νευρωνικό δίκτυο-Ιστορία

The term ‘neural network’ has its origins in attempts to find mathematical representations of information processing in biological systems (Widrow and Hoff, 1960; Rosenblatt, 1962; Rumelhart et al., 1986).

Indeed, it has been used very broadly to cover a wide range of different models, many of which have been the subject of exaggerated claims regarding their biological plausibility. From the perspective of practical applications of pattern recognition, however, biological realism would impose entirely unnecessary constraints.

Our focus is on neural networks as efficient models for statistical pattern recognition. In particular, we shall restrict our attention to the specific class of neural networks that have proven to be of greatest practical value, namely the multilayer perceptron.

[Pattern Recognition and Machine Learning, Bishop, Προσαρμογή από Κεφ. 5]

ALVINN '93, [Machine Learning, Mitchell]

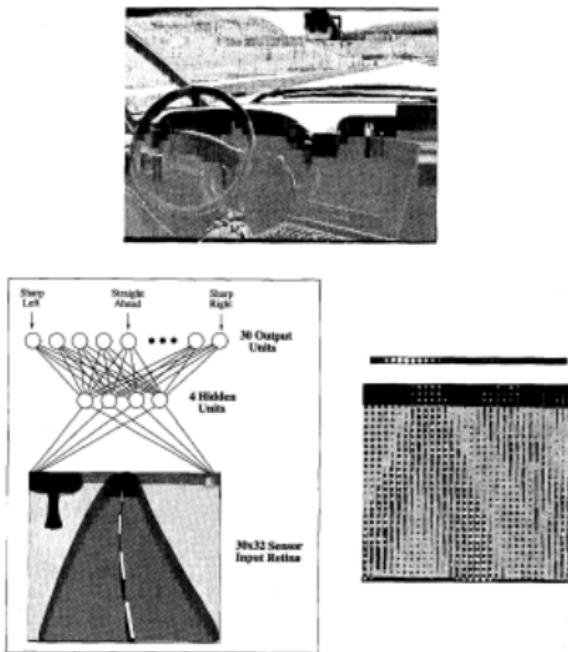


FIGURE 4.1

Neural network learning to steer an autonomous vehicle. The ALVINN system uses BACKPROPAGATION to learn to steer an autonomous vehicle (photo at top) driving at speeds up to 70 miles per hour. The diagram on the left shows how the image of a forward-mounted camera is mapped to 960 neural network inputs, which are fed forward to 4 hidden units, connected to 30 output units. Network outputs encode the commanded steering direction. The figure on the right shows weight values for one of the hidden units in this network. The 30×32 weights into the hidden unit are displayed in the large matrix, with white blocks indicating positive and black indicating negative weights. The weights from this hidden unit to the 30 output units are depicted by the smaller rectangular block directly above the large block. As can be seen from these output weights, activation of this particular hidden unit encourages a turn toward the left.

What is Deep Learning?

ARTIFICIAL INTELLIGENCE

Any technique that enables computers to mimic human behavior



MACHINE LEARNING

Ability to learn without explicitly being programmed

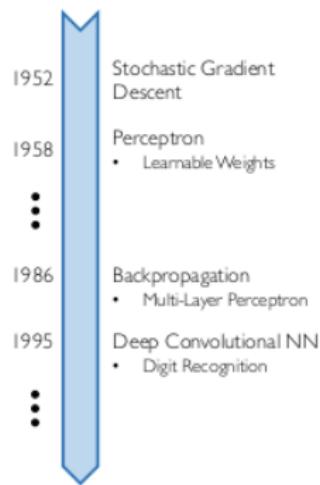


DEEP LEARNING

Extract patterns from data using neural networks



Why Now?



Neural Networks date back decades, so why the resurgence?

1. Big Data

- Larger Datasets
- Easier Collection & Storage



2. Hardware

- Graphics Processing Units (GPUs)
- Massively Parallelizable

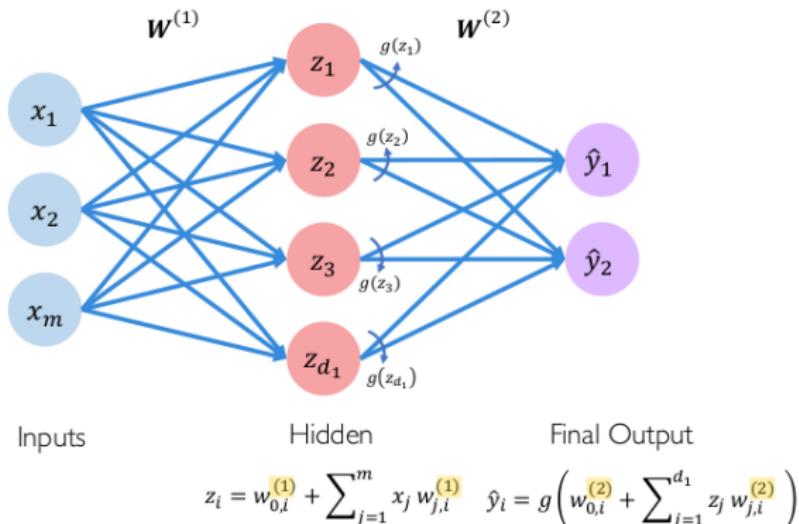


3. Software

- Improved Techniques
- New Models
- Toolboxes

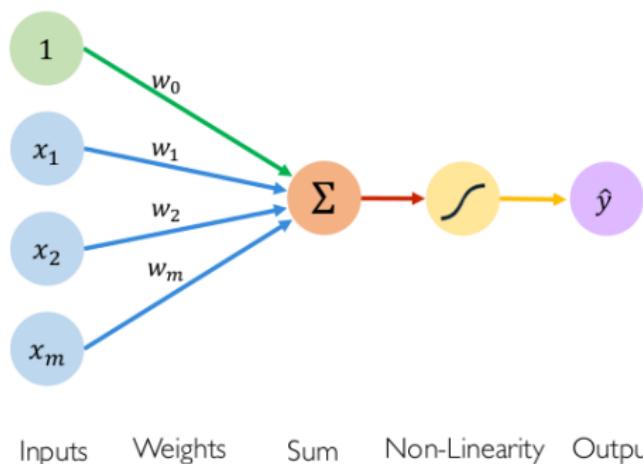


Single Layer Neural Network



NΔ-Forward Propagation

The Perceptron: Forward Propagation

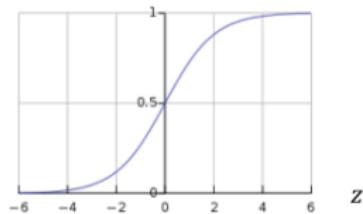


Activation Functions

$$\hat{y} = g(w_0 + X^T W)$$

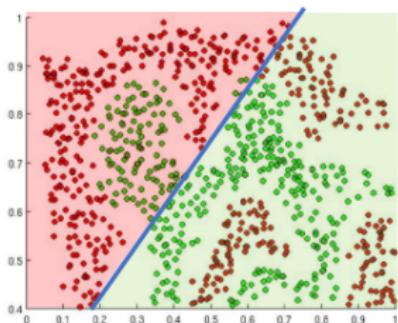
- Example: sigmoid function

$$g(z) = \sigma(z) = \frac{1}{1 + e^{-z}}$$

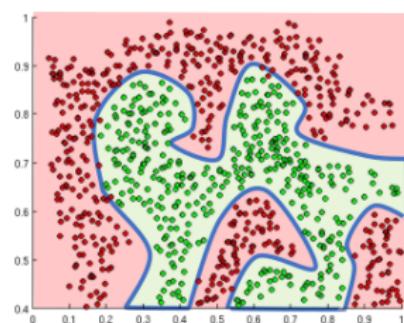


Importance of Activation Functions

The purpose of activation functions is to **introduce non-linearities** into the network



Linear Activation functions produce linear decisions no matter the network size



Non-linearities allow us to approximate arbitrarily complex functions

ΝΔ-Αρχιτεκτονική: Multilayer

2. Multilayer Feedforward Networks

The second class of a feedforward neural network distinguishes itself by the presence of one or more *hidden layers*, whose computation nodes are correspondingly called *hidden neurons* or *hidden units*. The function of hidden neurons is to intervene between the external input and the network output in some useful manner. By adding one or

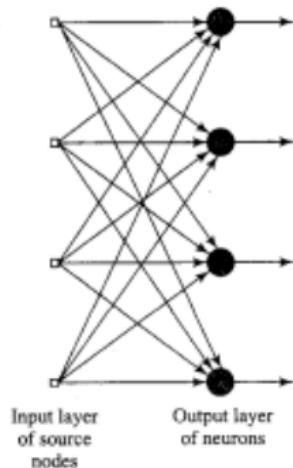


FIGURE 1.15 Feedforward or acyclic network with a single layer of neurons.

ΝΔ-Αρχιτεκτονική: Hidden

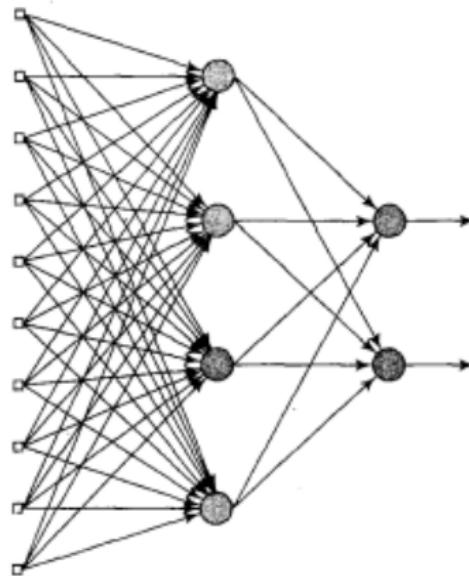


FIGURE 1.16 Fully connected feedforward or acyclic network with one hidden layer and one output layer.

Input layer
of source
nodes

Layer of
hidden
neurons

Layer of
output
neurons

ΝΔ-Αρχιτεκτονική: Recurrent

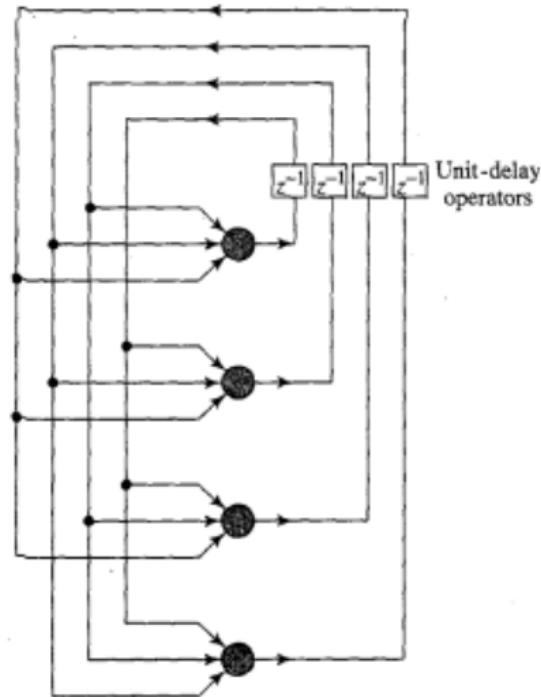


FIGURE 1.17 Recurrent network with no self-feedback loops and no hidden neurons.

ΝΔ-Αρχιτεκτονική: Combined

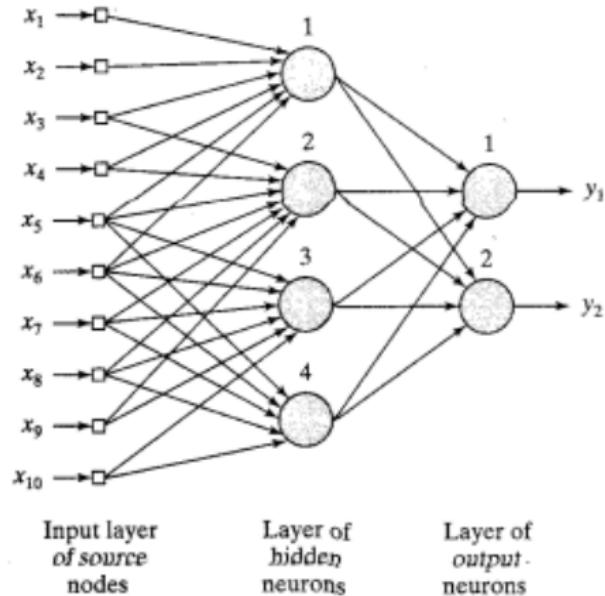


FIGURE 1.20 Illustrating the combined use of a receptive field and weight-sharing. All four hidden neurons share the same set of weights for their synaptic connections.

NΔ-Learning

2.1 INTRODUCTION

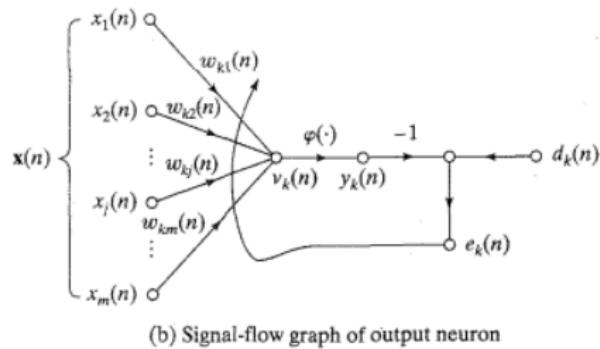
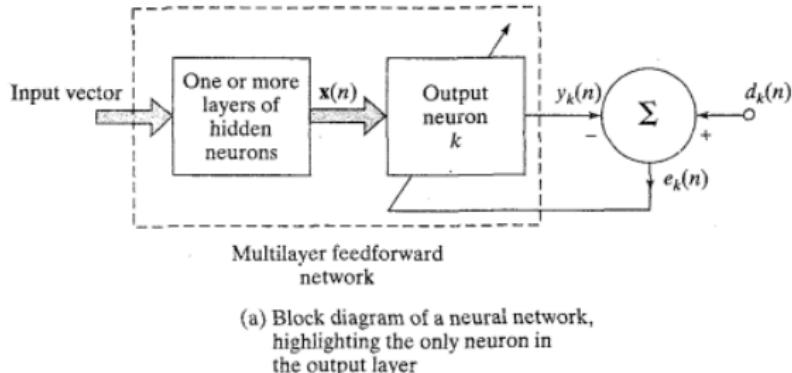
The property that is of primary significance for a neural network is the ability of the network to *learn* from its environment, and to *improve* its performance through learning. The improvement in performance takes place over time in accordance with some prescribed measure. A neural network learns about its environment through an interactive process of adjustments applied to its synaptic weights and bias levels. Ideally, the network becomes more knowledgeable about its environment after each iteration of the learning process.

There are too many activities associated with the notion of “learning” to justify defining it in a precise manner. Moreover, the process of learning is a matter of viewpoint, which makes it all the more difficult to agree on a precise definition of the term. For example, learning as viewed by a psychologist is quite different from learning in a classroom sense. Recognizing that our particular interest is in neural networks, we use a definition of learning that is adapted from Mendel and McLaren (1970).

We define learning in the context of neural networks as:

Learning is a process by which the free parameters of a neural network are adapted through a process of stimulation by the environment in which the network is embedded. The type of learning is determined by the manner in which the parameter changes take place.

NΔ-Learning



(b) Signal-flow graph of output neuron

FIGURE 2.1 Illustrating error-correction learning.

NΔ-Learning with Teacher

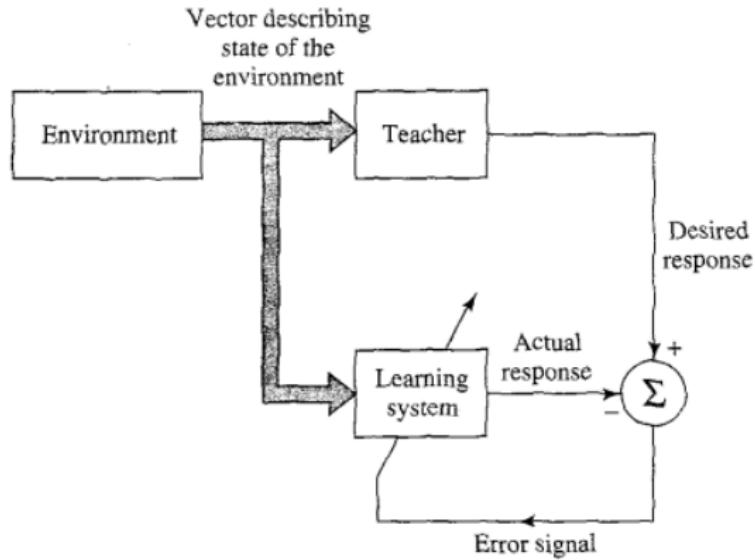


FIGURE 2.6 Block diagram of learning with a teacher.

NL Classification, [Machine Learning, Mitchell]

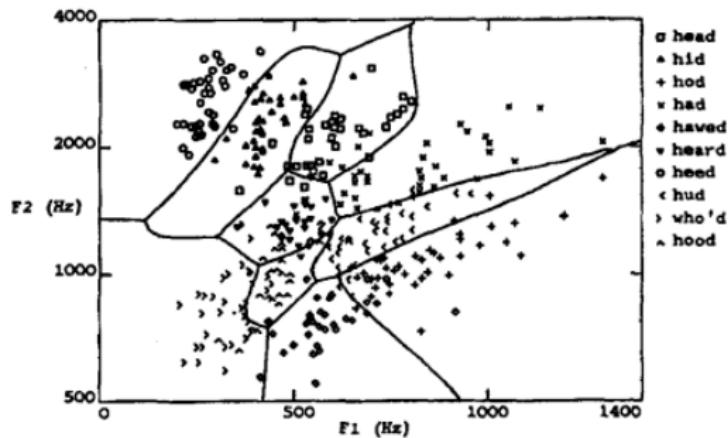
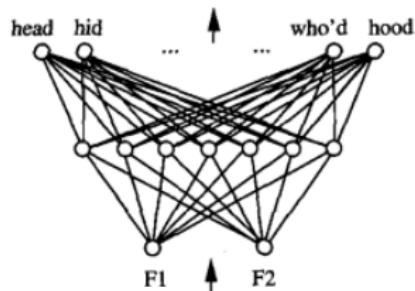


FIGURE 4.5

Decision regions of a multilayer feedforward network. The network shown here was trained to recognize 1 of 10 vowel sounds occurring in the context "h_d" (e.g., "had," "hid"). The network input consists of two parameters, F1 and F2, obtained from a spectral analysis of the sound. The 10 network outputs correspond to the 10 possible vowel sounds. The network prediction is the output whose value is highest. The plot on the right illustrates the highly nonlinear decision surface represented by the learned network. Points shown on the plot are test examples distinct from the examples used to train the network. (Reprinted by permission from Haung and Lippmann (1988).)

Ένας Virtual Sensor

Development of recurrent neural networks for virtual sensing of NOx emissions in internal combustion engines

Ivan Arsic, Cesare Pianese, Marco Sorrentino
Department of Mechanical Engineering, University of Salerno

Copyright © 2009 Society of Automotive Engineers, Inc.

ABSTRACT

The paper focuses on the experimental identification and validation of recurrent neural networks (RNN) for virtual sensing of NOx emissions in internal combustion engines (ICE). In particular, two different RNN architectures are proposed to improve RNN precision and generalization in predicting NO formation dynamics. The reference Spark Ignition (SI) engine was tested by means of an integrated system of hardware and software tools for model validation, control strategies prototyping. A fast response analyzer was used to measure NO emissions at the exhaust valve. The accuracy of the developed RNN model is assessed by comparing measured and predicted values for a wide range of operating scenarios. The results evidence that RNN-based virtual NO sensor will offer significant opportunities for implementing on-board feedforward and feedback control strategies aimed at improving the performance of after-treatment devices.

INTRODUCTION

Automotive engines and control systems are more and more sophisticated due to increasingly restrictive environmental regulations. Particularly in Diesel and GDI engines, complex after-treatment devices, such as SCR and DPF, have been introduced to meet the imposed NOx and particulate matter limits. Furthermore, on-board diagnostics for both Diesel and spark ignition (SI) engines is becoming tighter and tighter; in fact engine faults could lead to performance degradation, which may cause high fuel consumption. More recently, multi-fuel powertrain control units are required to identify the actual fuel composition and to adapt the strategies accordingly.

To cope with the new issues associated with novel complex hardware and to improve both powertrain performance and after-treatment efficiency, engine design methods and control and diagnostic strategies need to be refined. A key role in this context is played by performance as is played by on-board implementation of dedicated models, either as predictors or embedded in

model-based controllers or in diagnostics schemes. Virtual sensors (VSs) can be used as substitute of real sensors providing useful information about the actual process. In this scenario, virtual sensors are an interesting opportunity in achieving more challenging control and diagnostics targets.

In the last decade virtual sensors have been developed for many applications (e.g. aerospace, biotechnological processes, etc.) and their use has been validated in just few cases (Goodwin, 1999). VSs can be especially useful when: i) the direct measurements are impossible to perform, such as highly hazardous chemical or nuclear plants, (Savill and Puleo, 1998); ii) the measured variable is difficult to measure, requiring sensing characteristics (i.e. accuracy, dynamic performances); iii) the sensor is too expensive or iv) does not fit within the sensing location. It is also evident that VSs can be used for model validation and generation, and their use provides opportunities to implement innovative control schemes and diagnosis. Actually, VSs must be designed in such a way to simulate time-dependent processes and guarantee acceptable accuracy and short computation times. In this work it will be shown that the proposed virtual sensor also allows a limited recourse to experiments, thus reducing costs and developing time.

Internal combustion engines are complex systems to be controlled, with great precision and monitored continuously; moreover, their dynamic behavior requires both high sampling rates and rapid actuation. Therefore, in most industrial applications (e.g. model-based controllers, monitoring and fault diagnosis), the model must be fast enough to ensure the required performance. VSs may support the development of new control and diagnostic concepts in achieving more challenging targets. However, the development of a model is one of the most complex tasks to be accomplished in an engine; this is mainly due to the complexity of simulating a process that entails advanced thermo-chemistry and thermodynamics. In this context, the current state-of-the-art mainly relies on black-box applications where the process knowledge is stored in a large experimental database implemented on maps. A successful attempt to develop a steady-state first

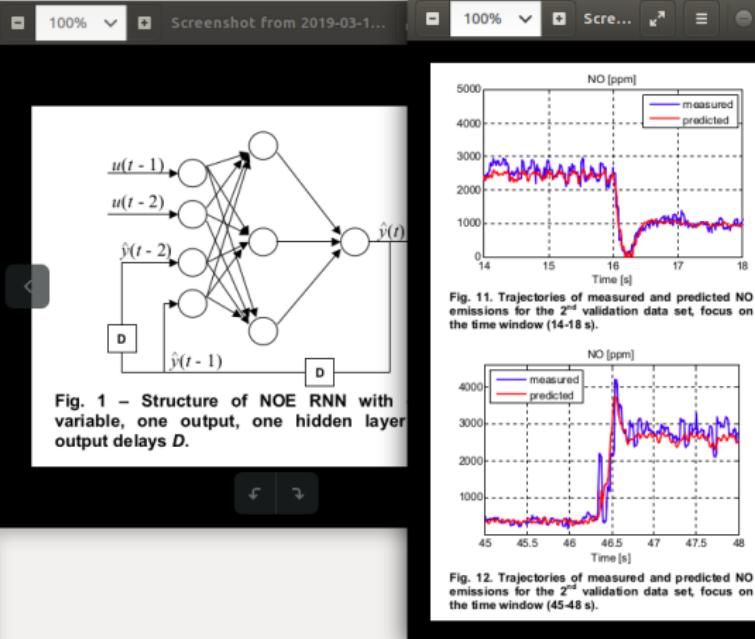


Fig. 1 – Structure of NOE RNN with variable, one output, one hidden layer output delays D .

Fig. 11. Trajectories of measured and predicted NO emissions for the 2nd validation data set, focus on the time window (14-18 s).

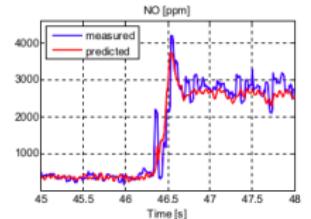
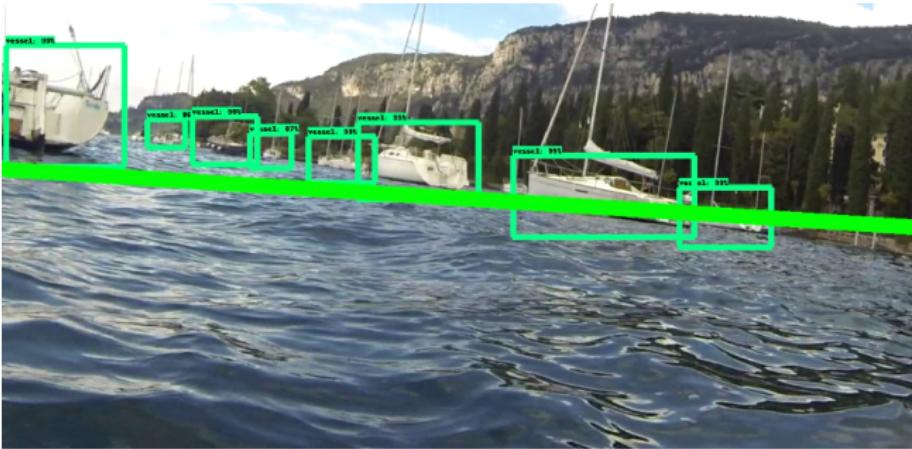


Fig. 12. Trajectories of measured and predicted NO emissions for the 2nd validation data set, focus on the time window (45-48 s).

ΝΔ σε θαλάσσιο περιβάλλον



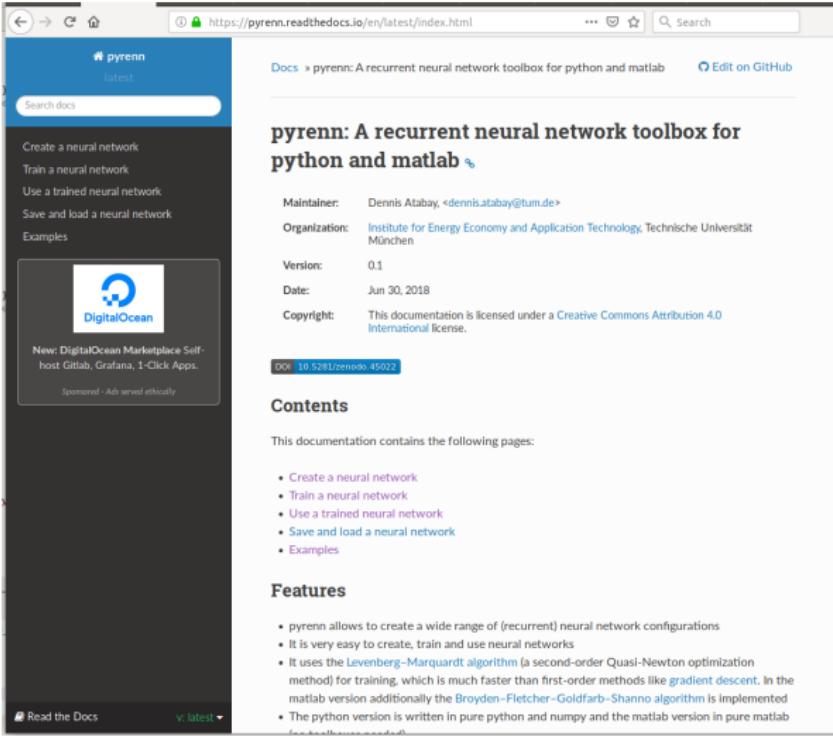
Διπλωματική Εργασία Κ. Βασιλόπουλου, ΣΝΜΜ 2018.

Σύγχρονα εργαλεία σε Python

- Βιβλίο Sebastian Raschka: Python Machine Learning, 2nd Edition.
- Γίνεται χρήση Python και βιβλιοθηκών numpy, SciPy, Scikit-learn για ΝΔ single - και multi - layer.
- <https://sebastianraschka.com/books.html>

PYRENN

- Ανοιχτή πλατφόρμα για σχεδίαση ΝΔ, με Python, MATLAB.



The screenshot shows a web browser displaying the documentation for the `pyrenn` toolbox. The URL is `https://pyrenn.readthedocs.io/en/latest/index.html`. The page title is "pyrenn: A recurrent neural network toolbox for python and matlab". It includes details about the maintainer (Dennis Atabay), organization (Institute for Energy Economy and Application Technology, Technische Universität München), version (0.1), date (Jun 30, 2018), and copyright (Creative Commons Attribution 4.0 International license). A DigitalOcean advertisement is visible on the left side. The main content area features sections for "Contents" and "Features", each with a bulleted list of topics or benefits. The bottom of the page has navigation icons for a presentation slide.

pyrenn
latest

Search docs

Create a neural network
Train a neural network
Use a trained neural network
Save and load a neural network
Examples

DigitalOcean

New: DigitalOcean Marketplace Self-host Gitlab, Grafana, 1-Click Apps.

Sponsored - Ads served ethically

DOI: 10.5281/zenodo.45022

Contents

This documentation contains the following pages:

- Create a neural network
- Train a neural network
- Use a trained neural network
- Save and load a neural network
- Examples

Features

- pyrenn allows to create a wide range of (recurrent) neural network configurations
- It is very easy to create, train and use neural networks
- It uses the Levenberg–Marquardt algorithm (a second-order Quasi-Newton optimization method) for training, which is much faster than first-order methods like gradient descent. In the matlab version additionally the Broyden–Fletcher–Goldfarb–Shanno algorithm is implemented
- The python version is written in pure python and numpy and the matlab version in pure matlab

Read the Docs v: latest

Tensor Flow

- TensorFlow is a free and open-source software library for dataflow and differentiable programming across a range of tasks. It is a symbolic math library, and is also used for machine learning applications such as neural networks. It is used for both research and production at Google...(Wikipedia)

The screenshot shows the TensorFlow.org homepage with the URL https://www.tensorflow.org/tutorials/get_started/intro_to_tensorflow. The page title is "Get Started with TensorFlow". The main content area displays a code snippet for a basic classification model:

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

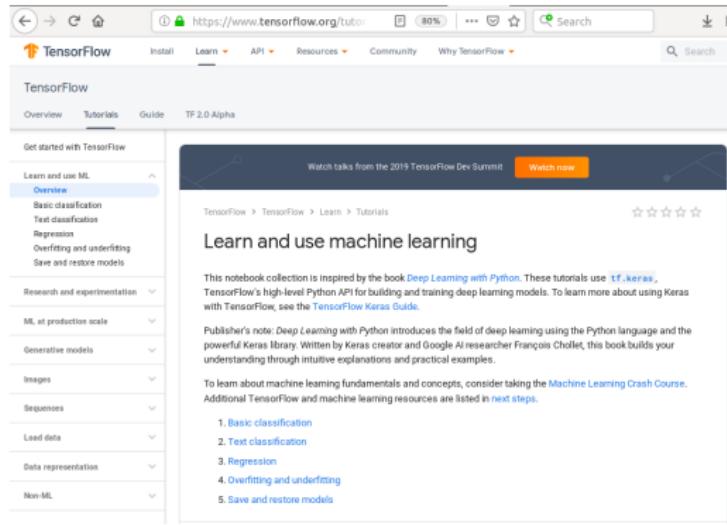
Below the code, there are two buttons: "Run code" and "View on GitHub".

The left sidebar contains navigation links for "Overview", "Tutorials", and "Guide", along with sections like "Learn and use ML", "Research and experimentation", and "ML at production scale".

The right sidebar includes sections for "ML at production scale" and "ML at research scale".

Tensor Flow

- Το βασικό interface είναι σε γλώσσα Python, αλλά για λόγους ταχύτητας είναι γραμμένο σε C++.



Tensor Flow

- Συνήθως το TF τρέχει ως η βασική μηχανή και σε υψηλότερο επίπεδο υπάρχουν (απλούστερα) APIs όπως το Keras κα.

The screenshot shows the TensorFlow.org website's 'Overview' page. At the top, there's a search bar and a navigation menu with links for 'Install', 'Learn', 'API', 'Resources', 'Community', and 'Why TensorFlow'. Below the menu, a main heading says 'TensorFlow makes it easy for beginners and experts to create machine learning models. See the sections below to get started.' There are two prominent buttons: 'See tutorials' and 'See the guide'. To the right, there's a 3D diagram of a server rack with orange lines representing data flow. Below the diagram, there are two sections: 'For beginners' and 'For experts'. The 'For beginners' section describes the Sequential API and provides a 'Hello World' example. The 'For experts' section describes the Subclassing API. At the bottom, there are two code snippets. The left snippet is for a simple neural network using the Sequential API:

```
import tensorflow as tf
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(input_shape=(28, 28)),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

The right snippet is for the Subclassing API:

```
class MyModel(tf.keras.Model):
    def __init__(self):
        super(MyModel, self).__init__()
        self.conv1 = Conv2D(32, 3, activation='relu')
        self.flatten = Flatten()
        self.d1 = Dense(128, activation='relu')
        self.d2 = Dense(10, activation='softmax')

    def call(self, x):
        x = self.conv1(x)
        x = self.flatten(x)
        x = self.d1(x)
        return self.d2(x)

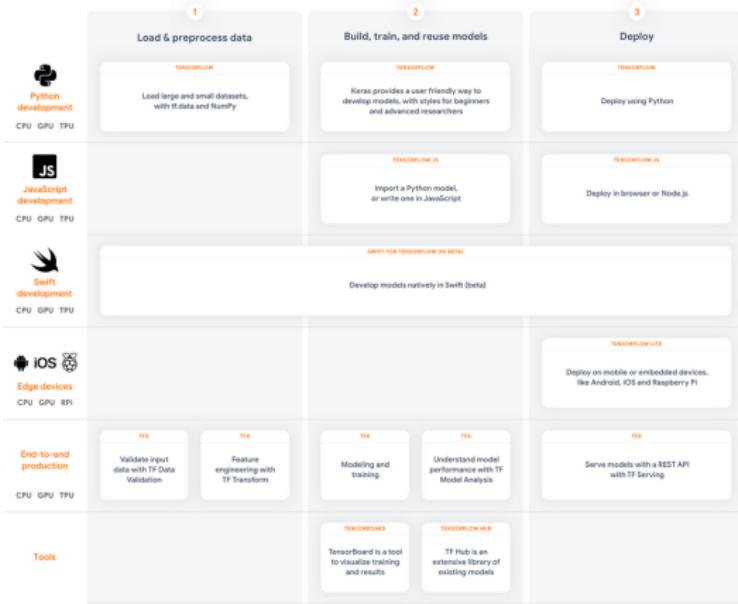
model = MyModel()

with tf.GradientTape() as tape:
    logits = model(images)
    loss_value = loss(logits, labels)
grads = tape.gradient(loss_value, model.trainable_variables)
optimizer.apply_gradients(zip(grads, model.trainable_variables))
```

At the very bottom, there are two buttons: 'Run code now!' and 'Try in Google's interactive notebook'.

Tensor Flow

- Το TF παρέχει μια συλλογή workflows για την ανάπτυξη και εκπαίδευση μοντέλων που χρησιμοποιούν Python, JavaScript για εύκολη ανάπτυξη σε cloud, browser ή συσκευές (πχ RPi), ανεξάρτητα από τη γλώσσα που χρησιμοποιείτε.



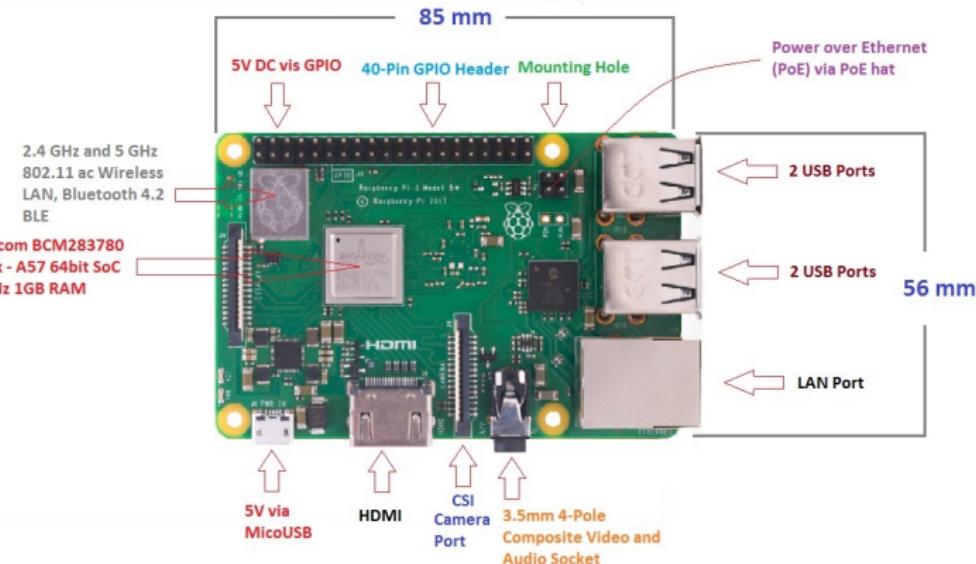
Tensor Flow - Coral

Coral Dev Board

CPU	i.MX 8M SoC w/ Quad-core A53
GPU	Integrated GC7000 Lite GPU
TPU	Google Edge TPU
RAM Memory	1GB LPDDR4 RAM
Flash Memory	8 GB eMMC
Security/Crypto	eMMC secure block for TrustZone MCHP ATECC608A Crypto Chip
Power	5V 3A via Type-C connector
Connectors	USB-C, RJ45, 3.5mm TRRS, HDMI
Supported OS	Mendel Linux (Debian derivative) Android
Supported ML	TensorFlow Lite



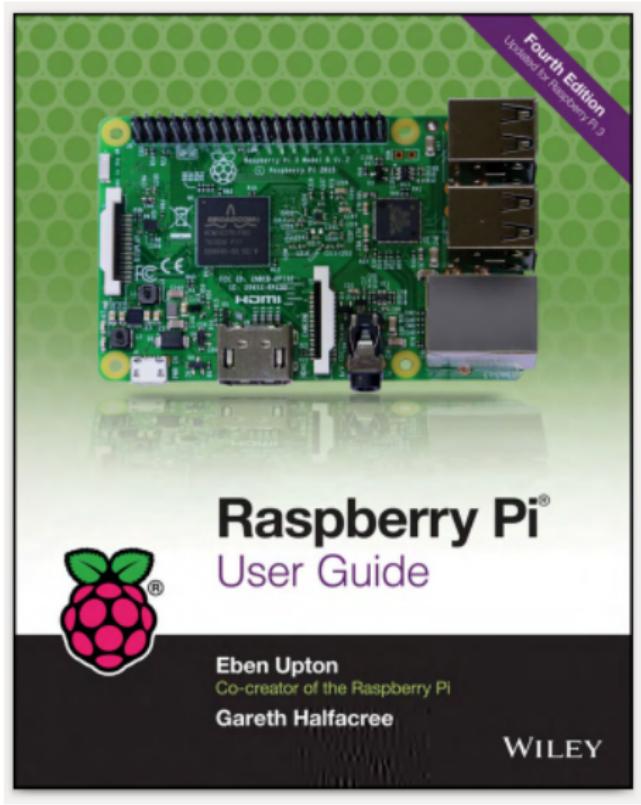
Raspberry Pi



Introduction to Raspberry Pi 3 B+

www.TheEngineeringProjects.com

Βιβλία RPi: E. Upton



Βιβλία RPi: E. Upton

GPIO Pinout Diagrams

Each pin of the GPIO port has its own purpose, with several pins working together to form particular interfaces. Figure 14-1 shows the layout of the GPIO port.

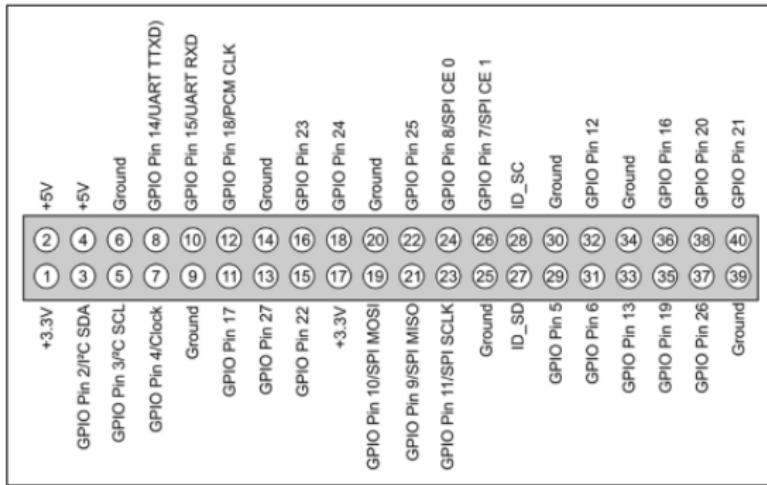


FIGURE 14-1: The Raspberry Pi's GPIO port and its pin definitions

Βιβλία RPi: E. Upton

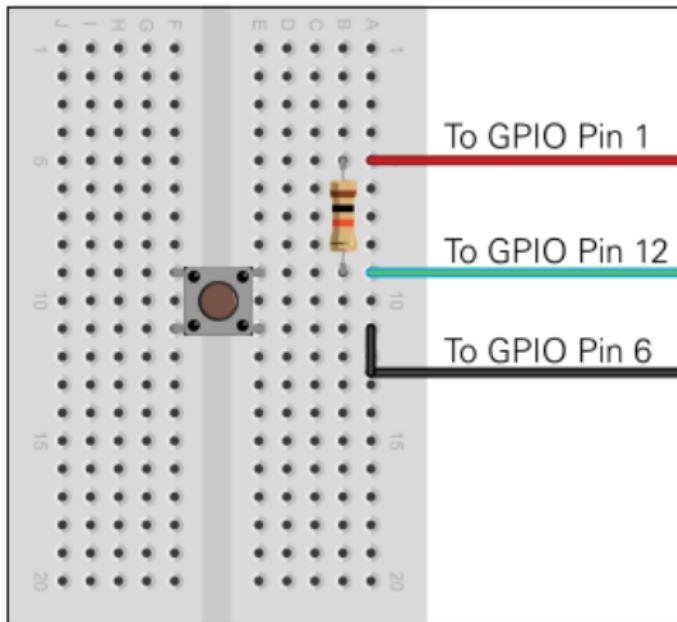


FIGURE 14-5: The example breadboard layout for a simple push-button input

Βιβλία RPi: E. Upton

CHAPTER 14 THE GPIO PORT

The final program should look like this:

```
import RPi.GPIO as GPIO
GPIO.setmode(GPIO.BOARD)
GPIO.setup(12, GPIO.IN)
while True:
    input_value = GPIO.input(12)
    if input_value == False:
        print("The button has been pressed.")
        while input_value == False:
            input_value = GPIO.input(12)
```

Βιβλία RPi: D. Molloy

'Exploring Raspberry Pi' is THE book to go to if you are interested in learning about the impressive physical computing capabilities of the Raspberry Pi platform. Derek Molloy imparts the electronics, programming, and embedded Linux skills that are vital to today's innovators in building the next generation of Internet of Things applications.*

Eben Upton, Co-creator of the Raspberry Pi

DEREK MOLLOY

EXPLORING RASPBERRY PI[®]

INTERFACING TO THE REAL WORLD WITH EMBEDDED LINUX™



WILEY

i2c bus: D. Molloy

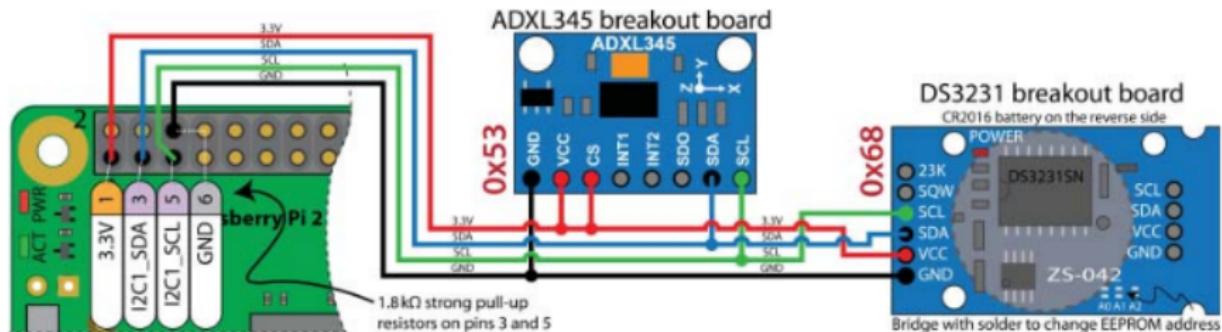


Figure 8-2: Two I²C devices connected to the I²C1 bus

SPI bus: D. Molloy

Connecting the ADXL345 to the RPi

The ADXL345 breakout board can be connected to the SPI bus as illustrated in Figure 8-12(a), where MOSI on the RPi is connected to SDA and MISO is connected to SDO. The clock lines and the slave select lines are also interconnected.

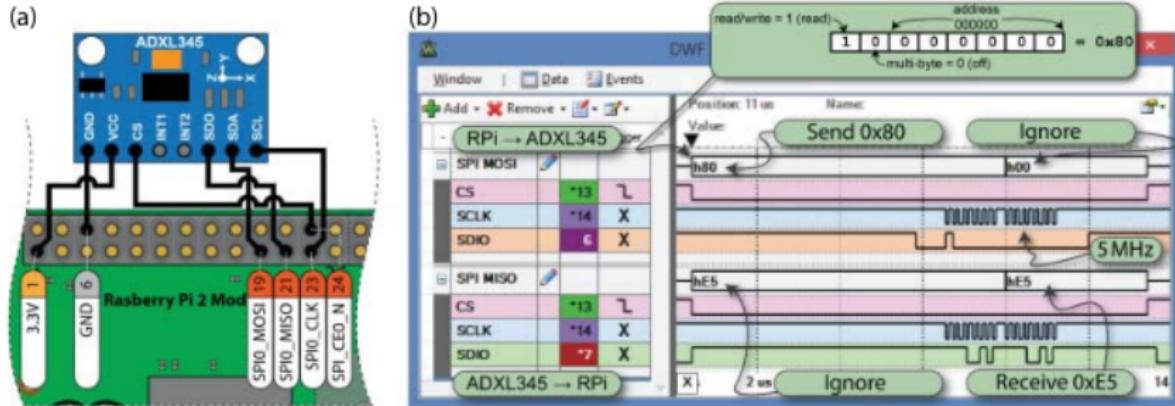


Figure 8-12: (a) SPI connection to the ADXL345; and (b) a capture of the communications required to read register 0x00